

Financial Statement Review Agent

2–3 minute live demo script · what to click · what to say

Field	Value
Document	DEMO-SCRIPT-001
Duration	2–3 minutes live (+ optional 30s Q&A setup)
Primary pack	Acme Demo Co (broken)
Backup pack	Acme Demo Co (clean)

1. Before you start (30 seconds silent prep)

1. UI open: `streamlit run ui/app.py`
2. Optional API tab: `http://localhost:8000/docs`
3. Sidebar set to **Acme Demo Co (broken)**, years 2023 / 2022
4. PDFs ready if asked: BRD · HLD · Phase2

2. Live script (speak this)

0:00–0:25 · **Problem**

“Financial statements still need manual review for math accuracy, year-on-year change, and consistency. That is slow and easy to miss. We built an AI-assisted review agent that checks the numbers with rules, explains findings with evidence, and can plug into any system through an API.”

0:25–0:45 · **How it works (one sentence architecture)**

“Rules own the numbers — totals and Balance Sheet identity. GenAI only writes the review wording from those results. It never invents amounts. The UI is one client; the same review runs over REST for ERP, chat, or batch.”

Point at sidebar Integration block for 2 seconds.

0:45–1:30 · **Live run (broken pack)**

Click: Company = broken → **Run review**

“On a broken pack we detect math failure on total assets and a Balance Sheet that does not balance — Assets not equal to Liabilities plus Equity — with expected versus actual evidence.”

Click: Findings tab → scroll to red MATH-TOTAL_ASSETS and CONS-01

Click: YoY variances → Revenue

Click: Observations → read one line → click Useful

“Year-on-year, revenue jumped about forty percent, so we flag it as material. Observations cite the check ids. Reviewer feedback is stored so we can improve wording over time.”

“Every run has a trace id. Completion and critical fails are notified — logged now, email or webhook when configured. Health endpoint tells ops if the service is down. Under load we can queue reviews; demo path is sync for speed.”

Optional click: Notifications tab · or browser /health · or Swagger POST /reviews

“We followed BRD → HLD → build with MVC services, tests, and CI. This is assistive review — not a statutory audit opinion — but it cuts review time and leaves an evidence trail. Happy to walk the architecture or the clean pack next.”

3. If judges ask

Question	Answer
Why DEGRADED?	No OpenAI key — templates still cite evidence. With a key, GenAI wording turns on.
Is this an audit?	No. Assistive review tool only.
Real data?	Demo fixtures now; Kaggle / client CSV maps to the same line codes.
Scale?	API + queue workers; UI stays a thin client.
Why not pure LLM?	Hallucinated totals are unsafe. Rules first, language second.

4. Slide order (5–7 slides)

- 1. Problem — manual FS review is slow / error-prone
- 2. Solution — rules + evidence + GenAI observations
- 3. Architecture — MVC · services · API-first (from HLD)
- 4. Live demo — broken pack screenshots or screen share
- 5. Platform — health · notify · feedback · integrate
- 6. SDLC — BRD → HLD → code → tests → CI
- 7. Roadmap — Kaggle breadth · deploy · Teams/Slack webhook